

## Hotel Reservation System

The screenshot displays the Expedia website interface for hotel reservations. At the top, there's a navigation bar with links like 'Home', 'Vacation Packages', 'Hotels', 'Cars', 'Flights', 'Cruises', 'Things to Do', 'DEALS & OFFERS', 'GROUPON Getaways', and 'Rewards'. A 'Hotwire' banner offers '4-star hotels, 2-star prices' with a 'SHOW ME' button. The main search results are for 'Hotels in New York (and vicinity) on Apr 10 – Apr 11 for 2 adults', showing 514 results. A summary table indicates average prices: Hotel avg. \$226, 3 star avg. \$193, and 4 star avg. \$254. Below this, three hotel listings are shown: Hampton Inn New York Chelsea (4.4 out of 5, \$249/night), The Belvedere Hotel (4.3 out of 5, \$189/night), and Hilton New York Fashion District (4.4 out of 5, \$189/night). A sidebar on the left provides filters for 'Hotel Name Contains', 'Star Rating', 'Price (Avg. per night)', 'View Hotels On a Map', and 'Neighborhood'. On the right, there are sponsored listings for NYU Stern Executive MBA and various NYC hotels.

## High Level Features

- User browses through rooms available for given date range
- User reserves a type of room in a particular hotel
- One check-in, hotel manager assigns a room of that type to the user

## Performance Considerations

- When does WRITE happen?
  - User reserves a room
  - User cancels a reservation
  - New hotel or room added
- When does READ happen?

- Browsing through hotel catalog
- Browsing through hotel features

So a significantly higher amount of READ than WRITE

## API Requirement

Generic CRUD endpoints for hotel and room management. Let's ignore them.

## Reservation

GET /reservations

GET /reservations/123

POST /reservations

DELETE /reservations/123

## Data Model

Let's go with a relational database like MySQL or PostgreSQL.

Why?

- Easier to model hotel and reservation data
- More READ than WRITE
- Mostly CRUD operations
- ACID properties, transactional guarantees
- Easier locking mechanisms
- Data can be easily sharded for scalability

## Hotel Table

hotel_id	name	address	location
----------	------	---------	----------

## Room Table

room_id	room_type_id	hotel_id	is_available
---------	--------------	----------	--------------

## Room Type Inventory Table

hotel_id	room_type_id	date	total_inventory	total_reserved

## Rate Table

hotel_id	room_type_id	date	rate
----------	--------------	------	------

## Reservation Table

reservation_id	hotel_id	room_type_id	start_date	end_date	status	guest_id
----------------	----------	--------------	------------	----------	--------	----------

## Guest Table

guest_id	first_name	last_name	age
----------	------------	-----------	-----

## What happens when user wants to reserve?

- Check the "Room Type Inventory" table for availability
- If not available:
  - Don't reserve. Throw error.
- If available:
  - Update inventory
  - Create reservation
  - **Both should be done in the same transaction**

## How to avoid double booking by the same user?

Let's say the user clicked "Book" twice in very quick succession.  
How do we avoid booking twice?

## Use an idempotency key

Steps:

- When user lands in final checkout page

- Backend generates a unique key (reservation\_id)
- Sends the key to the client
- Client sends the key to the API when reserving
- If user clicks twice:
  - Same key goes to the backend
  - Backend knows a reservation with that key has already been created
  - So Backend throws away the request

Simpler Solution:

- Just gray out and disable the button on client side after being clicked once
- Problem
  - User can disable JavaScript and get around it

## How to avoid multiple users reserving the same room?

### Approach 1: Use Locking

- Add a new column `version` to the tables
- Client reads the version column when reading a row
- When writing, the application increments the version by 1.
- In the meantime, if version has already been incremented by a different client:
  - Database throws an error
  - Operation is rolled back
  - User will have to try again with a different room

### Approach 2: Database Constraint (If supported)

- Add a database constraint
- `CHECK((total_inventory - total_reserved) >= 0)`
- If constraint fails when writing, transaction is rolled back

## How to scale?

### Database getting too large?

- Nightly batch to remove & archive older rows

- Shard database by `hotel_id`

## Read is taking too long?

- Move read traffic from database to cache
- For more popular hotels, cache the inventory information
  - Will lead to more user facing errors
  - Inconsistent inventory data between cache and database
- For all hotels, cache static data like features and hotel details

## How can you improve cache data accuracy?

- Database CDC updates Cache
- Whenever inventory changes, cache is invalidated and updated with new inventory